

JMM JIT GC



JAVA DEVELOPER

VS

JAVA DEVELOPER

JAVA DEVELOPER

The internet will make those bad words go away





Essential


Googling the Error Message

ONLY?

The Practical Developer
@ThePracticalDev



JAVA DEVELOPER




- **JMM**
- **JIT**
- **GC**

JAVA MEMORY MODEL

1980s – one CPU instruction execution == one RAM access action


2010s – CPU performance increase - ~10_000 times
RAM performance increase - ~10 times



1 CPU cycle	0.3 ns
Level 1 cache access	0.9 ns
Level 2 cache access	2.8 ns
Level 3 cache access	12.9 ns
Main memory access	120 ns
Solid-state disk I/O	50-150 μs
Rotational disk I/O	1-10 ms
Internet: SF to NYC	40 ms
Internet: SF to UK	81 ms
Internet: SF to Australia	183 ms


1980s – one CPU instruction execution == one RAM access action

2010s – CPU performance increase - ~10_000 times
RAM performance increase - ~10 times



1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid-state disk I/O	50-150 μs	2-6 days
Rotational disk I/O	1-10 ms	1-12 months
Internet: SF to NYC	40 ms	4 years
Internet: SF to UK	81 ms	8 years
Internet: SF to Australia	183 ms	19 years

HEAVY OPTIMIZATION NEEDED



1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access	120 ns	6 min
Solid-state disk I/O	50-150 μs	2-6 days
Rotational disk I/O	1-10 ms	1-12 months
Internet: SF to NYC	40 ms	4 years
Internet: SF to UK	81 ms	8 years
Internet: SF to Australia	183 ms	19 years



```
int foo = 0;  
int bar = 0;
```

```
void run {  
    foo += 1;  
    bar += 1;  
  
    foo += 2;  
}
```

```
int foo = 0;  
int bar = 0;
```

```
void run {  
    foo += 1;  
    bar += 1;  
  
    foo += 2;  
}
```

```
void run {  
    foo += 1;  
    foo += 2;  
  
    bar += 1;  
}
```




REORDERING

```
int foo = 0;  
int bar = 0;
```

```
void run {  
    foo += 1;  
    bar += 1;  
  
    foo += 2;  
}
```

```
void run {  
    foo += 1;  
    foo += 2;  
  
    bar += 1;  
}
```



REORDERING


foo == 0, bar == 0
foo == 1, bar == 0
foo == 1, bar == 1
foo == 3, bar == 1

foo == 0, bar == 0
foo == 1, bar == 0
foo == 3, bar == 0
foo == 3, bar == 1

```
int foo = 0;  
int bar = 0;
```

```
void run {  
    foo += 1;  
    bar += 1;  
  
    foo += 2;  
}
```

```
void run {  
    foo += 3;  
    bar += 1;  
}
```



REORDERING

foo == 0, bar == 0
foo == 1, bar == 0
foo == 1, bar == 1
foo == 3, bar == 1

foo == 0, bar == 0
???
foo == 3, bar == 0
foo == 3, bar == 1


```
boolean flag = true;  
int count = 0;
```

```
void run { //thread1  
    while (flag) {  
        count++;  
    }  
}
```


```
void run { //thread2  
    flag = false;  
}
```

```
boolean flag = true;  
int count = 0;
```

JIT OPTIMIZATION

```
void run { //thread1  
    while (flag) {  
        count++;  
    }  
}
```

```
void run { //thread1  
    while (true) {  
        count++;  
    }  
}
```




```
void run { //thread2  
    flag = false;  
}
```

```
boolean flag = true;  
int count = 0;
```

JIT OPTIMIZATION

```
void run { //thread1  
    while (flag) {  
        count++;  
    }  
}
```

```
void run { //thread1  
    while (true) {  
        count++;  
    }  
}
```



```
void run { //thread2  
    flag = false;  
}
```

```
void run { //thread2  
    //flag = false;  
}
```

JAVA **M**EMORY **M**ODEL

answers the question:

“What values can we observe upon reading from a specific field?”

JAVA MEMORY MODEL

answers the question:

“What values can we observe upon reading from a specific field?”

```
int foo = 0;

//single-threaded
void run {
    foo = 1;
    if (foo < 1) {
        throw new ISE();
    }
}
```

JAVA MEMORY MODEL

answers the question:

“What values can we observe upon reading from a specific field?”

```
int foo = 0;
```

```
//single-threaded
```

```
void run {
```

```
    foo = 1;
```

```
    if (foo < 1) {  
        throw new ISE();
```

```
    }
```

```
}
```

//->

1.

//->

2.



program order

JAVA MEMORY MODEL

answers the question:

“What values can we observe upon reading from a specific field?”

```
int foo = 0;
int bar = 0;
//single-threaded
void run {
    //bar = 1;
    foo = 1;
    //bar = 1;
    if (foo < 1) {
        throw new ISE();
    }
    //bar = 1;
}
```

JAVA MEMORY MODEL

volatile
synchronized
java.util.concurrent.*


```
int foo = 0;  
int bar = 0;  
int baz = 0;
```

```
//single-threaded  
void run {  
    foo += 1;  
    bar += 1;  
    baz += 1;  
}
```

volatile

```
int foo = 0;
int bar = 0;
int baz = 0;

//single-threaded
void run {
    foo += 1;
    bar += 1;
    baz += 1;
}
```

volatile

```
void run {
    baz += 1;
    foo += 1;
    bar += 1;
}

void run {
    foo += 1;
    baz += 1;
    bar += 1;
}

void run {
    bar += 1;
    foo += 1;
    baz += 1;
}
```

volatile

```
int foo = 0;  
volatile int bar = 0;  
int baz = 0;
```

```
//single-threaded
```

```
void run {  
    foo += 1;  
    bar += 1;  
    baz += 1;  
}
```

//-> PROGRAM ORDER BARRIER

```
int foo = 0;  
volatile int bar = 0;  
int baz = 0;
```

volatile

time

```
//multi-threaded  
void run { //thread1  
    foo += 1;  
    bar += 1;  
    baz += 1;  
}
```

happens-before

```
//multi-threaded  
void run { //thread2  
    foo += 1;  
    bar += 1;  
    baz += 1;  
}
```


JAVA MEMORY MODEL

“happens-before” order

- action in a thread *happens-before* any subsequent action in that thread
- writing to a volatile variable *happens-before* reading from same variable
- `thread.start()` *happens-before* any action in that thread
- all actions in a thread *happen-before* a return from a call to `.join()` on that thread
- monitor release *happens-before* any subsequent monitor acquisition on the same object

Copyrighted Material
BRIAN GOETZ



WITH TIM PEIERLS, JOSHUA BLOCH,
JOSEPH BOWBEER, DAVID HOLMES,
AND DOUG LEA

JAVA **CONCURRENCY** **IN PRACTICE**



Chapter 17



- **JMM**
- **JIT**
- **GC**

JUST-IN-TIME COMPILATION

AOT -> **Ahead-Of-Time compilation**
happens prior to execution

JIT -> **Just-In-Time compilation**
happens during execution

AOT

○ Pros

- **can perform time-consuming analysis and complex optimizations**

○ Cons

- **static information is not always enough to optimize effectively**
- **cannot use profiling**
- **cannot make use of previously unavailable hardware features**


JIT



○ Pros

- **uses aggressive speculative optimizations based on profiling information**
- **can use new hardware features out-of-the-box**

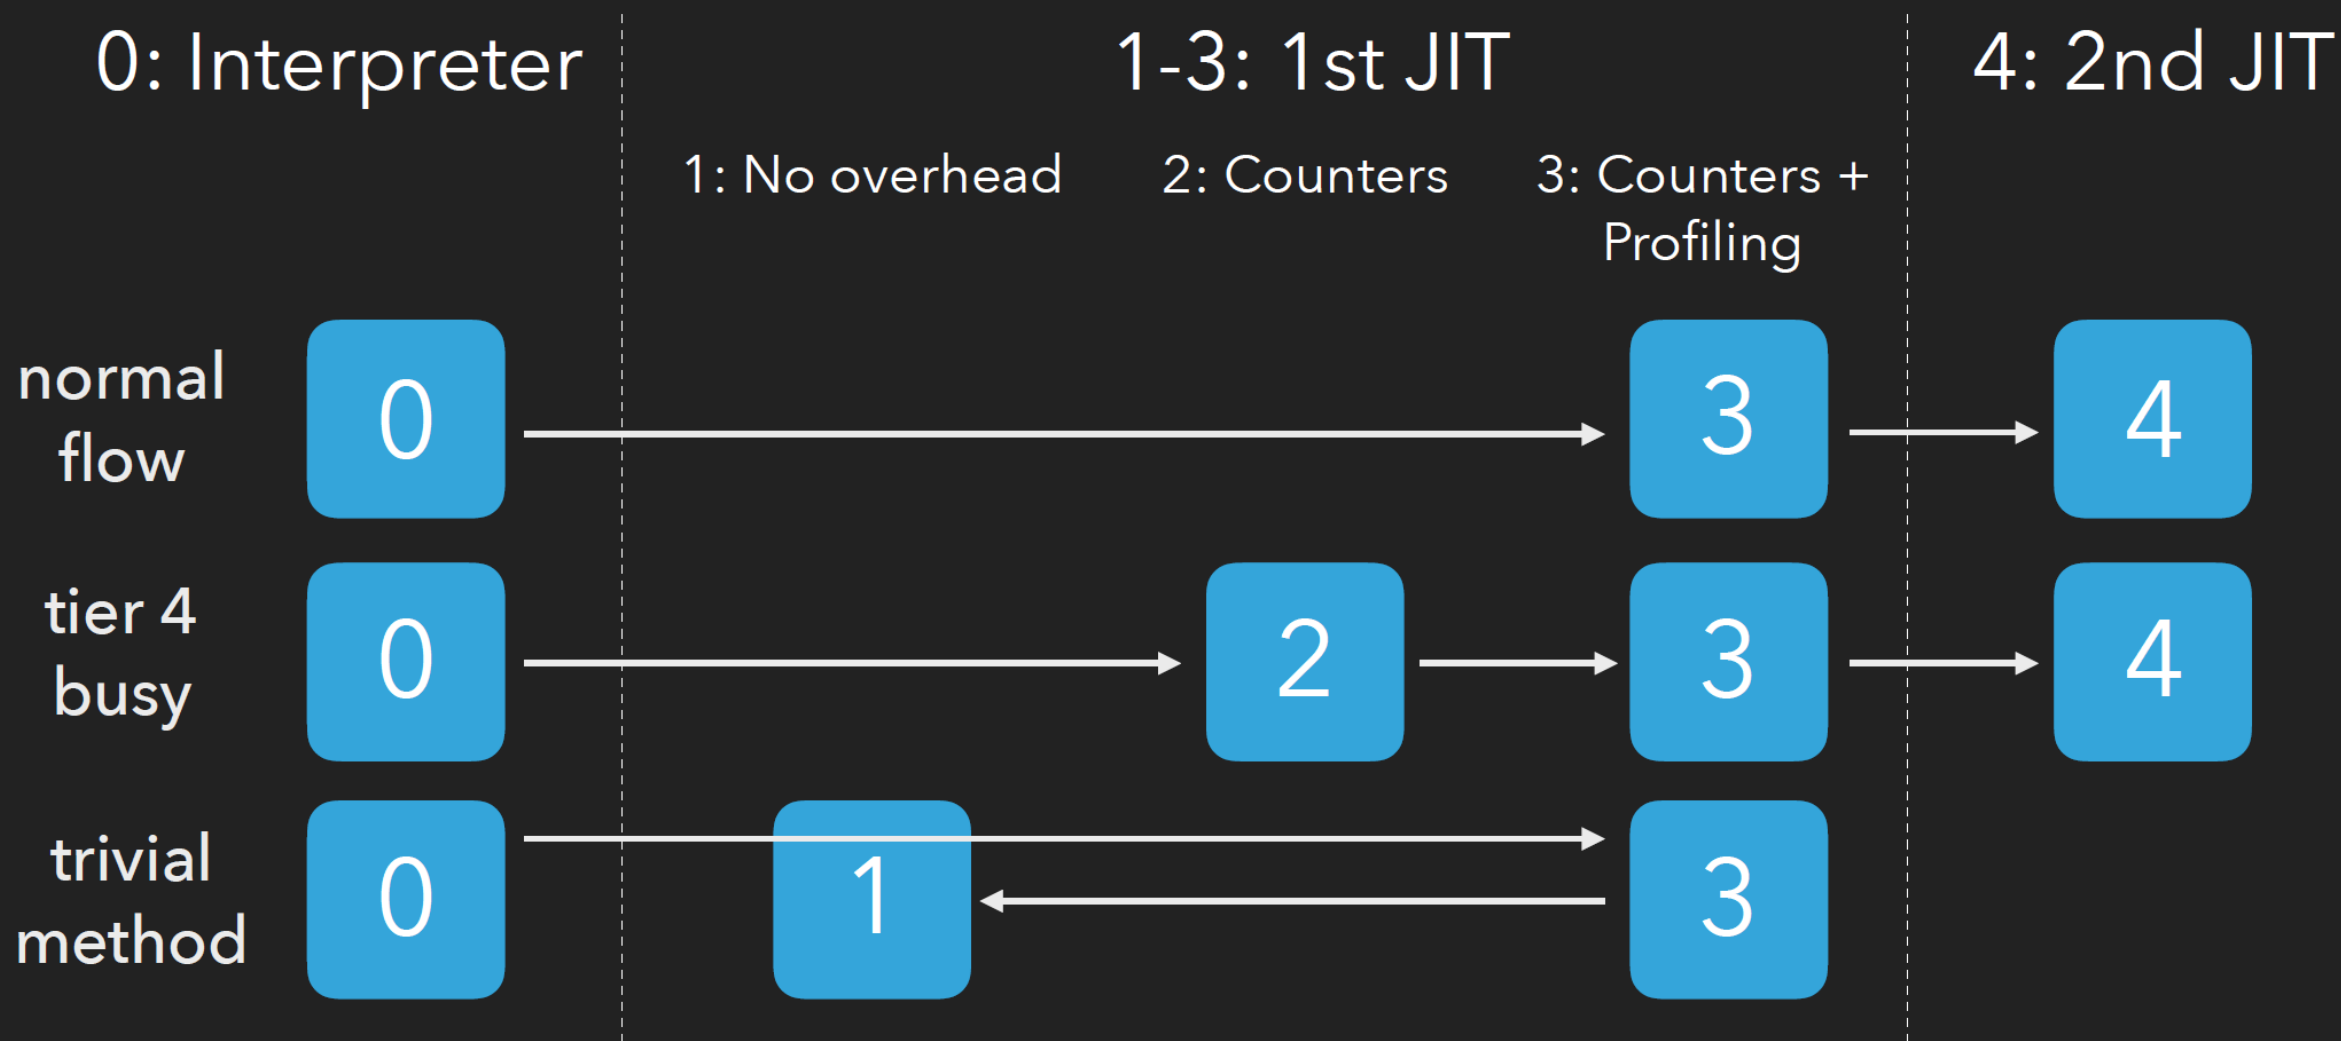
○ Cons

- **uses resources at runtime**
- **can negatively affect startup time**
- **peak performance can be lower**

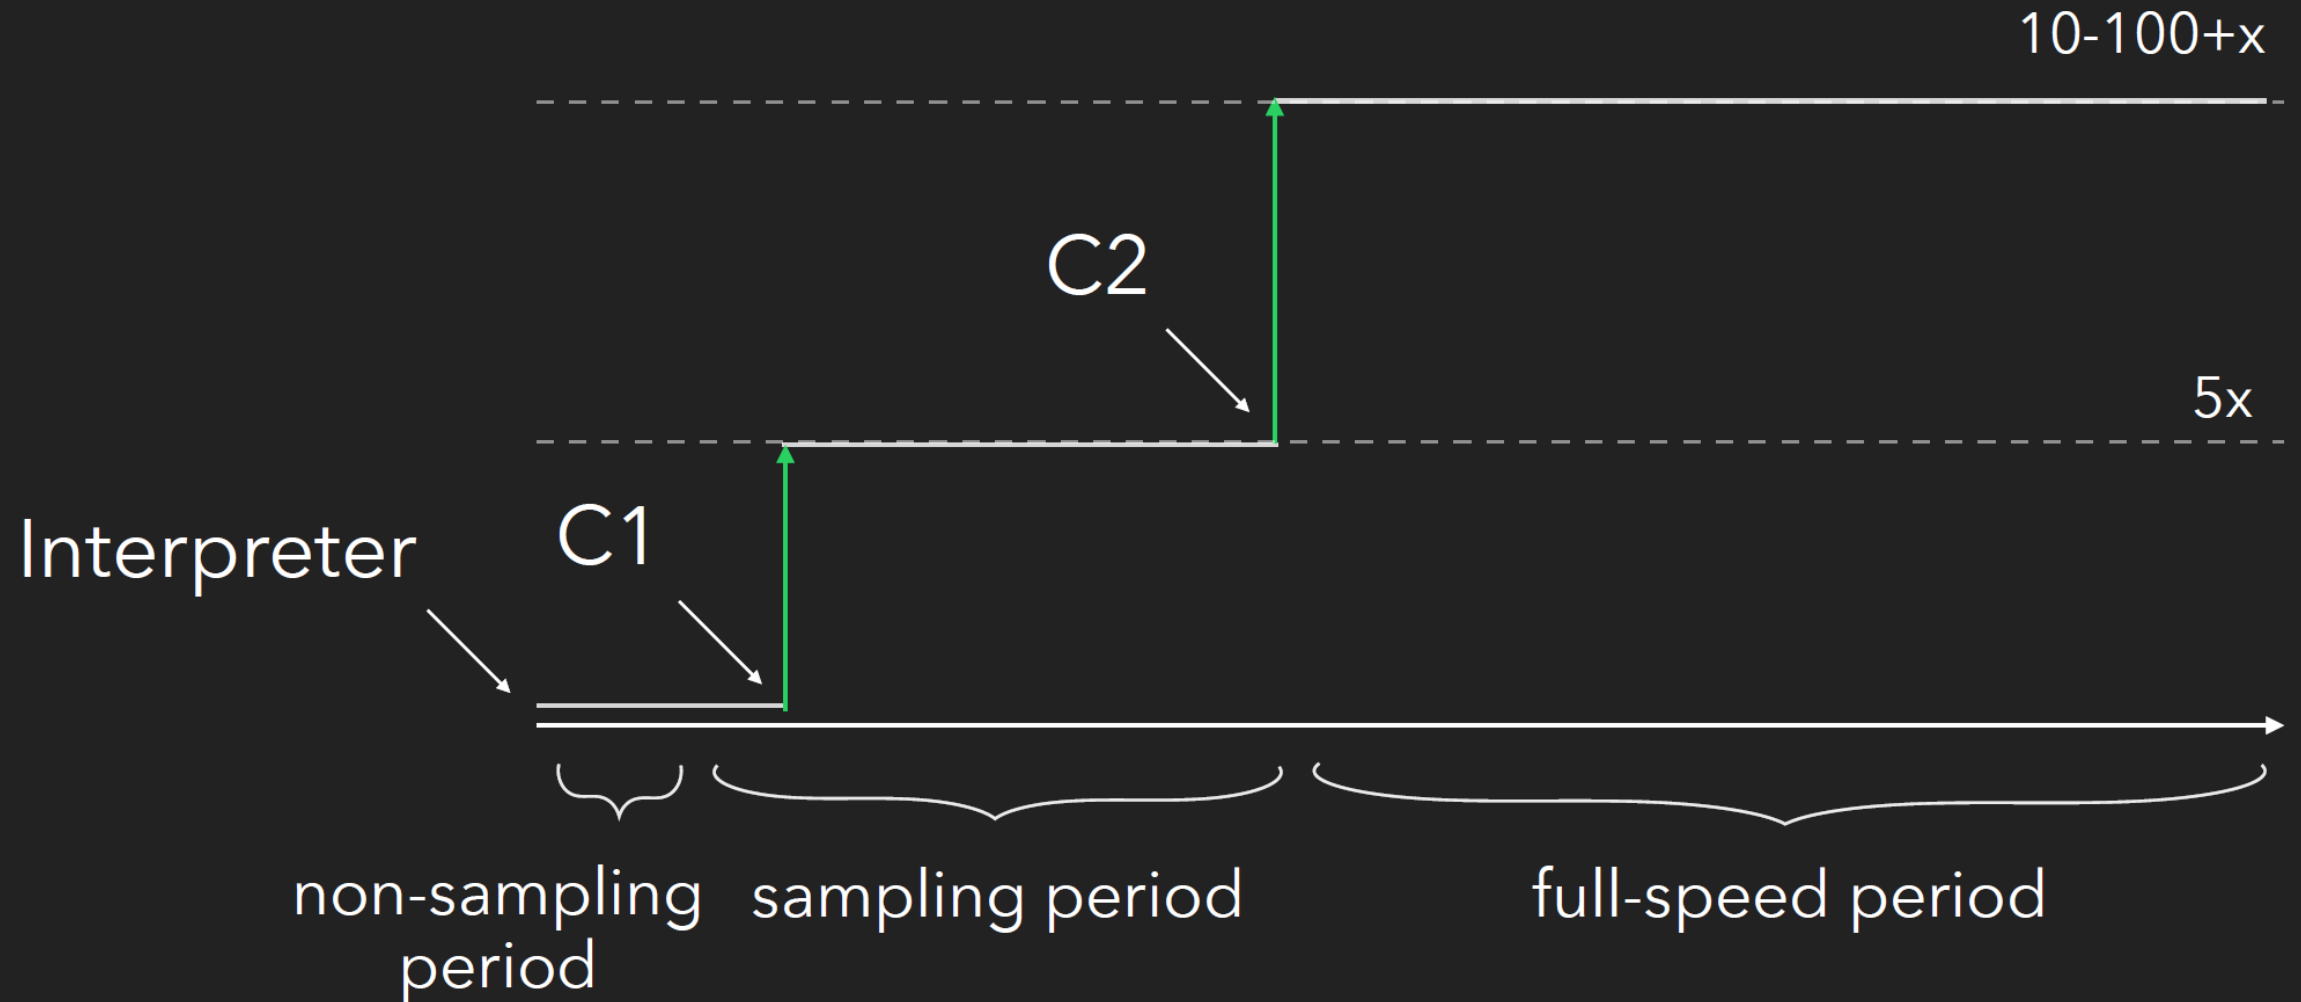




TIERED COMPIRATION



MENTAL MODEL



JIT OPTIMIZATIONS

compiler tactics

- delayed compilation
- tiered compilation
- on-stack replacement
- delayed reoptimization
- program dependence graph representation
- static single assignment representation

proof-based techniques

- exact type inference
- memory value inference
- memory value tracking
- constant folding
- reassociation
- operator strength reduction
- null check elimination
- type test strength reduction
- type test elimination
- algebraic simplification
- common subexpression elimination
- integer range typing

flow-sensitive rewrites

- conditional constant propagation
- dominating test detection
- flow-carried type narrowing
- dead code elimination

language-specific techniques

- class hierarchy analysis
- devirtualization
- symbolic constant propagation
- autobox elimination
- escape analysis
- lock elision
- lock fusion
- de-reflection

speculative (profile-based) techniques

- optimistic nullness assertions
- optimistic type assertions
- optimistic type strengthening
- optimistic array length strengthening
- untaken branch pruning
- optimistic N-morphic inlining
- branch frequency prediction
- call frequency prediction

memory and placement transformation

- expression hoisting
- expression sinking
- redundant store elimination
- adjacent store fusion
- card-mark elimination
- merge-point splitting

loop transformations

- loop unrolling
- loop peeling
- safepoint elimination
- iteration range splitting
- range check elimination
- loop vectorization

global code shaping

- inlining (graph integration)
- global code motion
- heat-based code layout
- switch balancing
- throw inlining

control flow graph transformation

- local code scheduling
- local code bundling
- delay slot filling
- graph-coloring register allocation
- linear scan register allocation
- live range splitting
- copy coalescing
- constant splitting
- copy removal
- address mode matching
- instruction peepholing
- DFA-based code generator

```
void run {  
    int number = 2;  
    for (int i = 0; i < 10; i++) {  
        number = addSelf(number);  
    }  
}
```

```
int addSelf(int num) {  
    return num + num;  
}
```



```
void run {  
    int number = 2;  
    for (int i = 0; i < 10; i++) {  
        number = number + number;  
    }  
}
```

OPENS WAY FOR COUNTLESS OTHER OPTIMIZATIONS



METHODS


- Monomorphic
- Bimorphic
- Polymorphic
- Megamorphic



DEOPTIMIZATION




REVISED MENTAL MODEL



-XX: +PrintCompilation

-XX: +UnlockDiagnosticVMOptions

-XX: +PrintInlining



- **JMM**
- **JIT**
- **GC**

GARBAGE COLLECTION

GARBAGE COLLECTION

parallel vs concurrent

concurrent vs stop-the-world

monolithic vs incremental

GARBAGE **C**OLLECTION

parallel vs concurrent


concurrent vs stop-the-world

monolithic vs incremental

(GC) SAFEPPOINTS

LIVE SET

MARK




SWEEP

Before



After
sweep



free-list


linear to the heap size

SWEEP


Before



After
sweep



java.lang.OutOfMemoryError




MARK/SWEEP/COMPACT

relocate -> remap


Before



After
sweep




After
compact



linear to the live set

(MARK +) COPY



linear to the live set; single pass

requires 2x live set memory

GARBAGE COLLECTION

mark-sweep-compact, copy, mark-compact, mark-sweep-(compact)

- ***COPY***
 - *needs 2x live set memory*
 - *linear to live set*
 - *Monolithic*
- ***MARK-COMPACT***
 - *needs 2x live set memory*
 - *linear to live set*
- ***MARK-SWEEP-COMPACT***
 - *1x live set memory*
 - *linear to heap size (during sweep)*
- ***MARK-SWEEP-(COMPACT)***
 - *compacts periodically*

GENERATIONAL COLLECTION


based on a 'Weak Generation' hypothesis

“MOST OBJECTS DIE YOUNG”

- **separate young objects from the older ones**
- **use a Copy collector (linear to the live set)**
- **promote object that are old enough to another area**
- **only collect older areas as needed**

➤ *requires* "remembered sets"


➤ *does NOT need 2x memory*



Minor GC



Full GC



Serial GC

-XX:+UseSerialGC

- **monolithic stop-the-world** **Copy** in Young Generation
- **monolithic stop-the-world** **Mark-Sweep-Compact** in Old Generation

Best for ~100MB heaps with single-core CPU

Parallel GC

default in Java 7 & 8

`-XX:+UseParallelGC` `-XX:+UseParallelOldGC`

- **monolithic stop-the-world** **Copy** in Young Generation
- **monolithic stop-the-world** **Mark-Sweep-Compact** in Old Generation

Best for multi-core CPU --> high throughput

“Mostly Concurrent Mark and Sweep Garbage Collector”

–XX: +UseConcMarkSweepGC

- **monolithic stop-the-world **Copy** in Young Generation in parallel**
- **mostly concurrent **Mark-Sweep** in Old Generation**

Best for *low latency*, uses application resources.

*Fallback for Full Collection (monolithic STW)
when compaction is needed*

G1 – “Garbage First”

default since Java 9

`-XX:+UseG1GC`

- **monolithic stop-the-world *Copy* in Young Generation in parallel**
- ***mostly* concurrent *Mark* in Old Generation**
- ***mostly* incremental stop-the-world *Compact* in Old Generation**

Best for specific goals
Tries to delay Full GC

